

Lesson 1: Introduction to Fault-Tolerant Computing

1. Introduction

In computing systems, failures can occur due to various reasons, such as hardware malfunctions, software bugs, or external factors like power outages. These failures can lead to system crashes, data loss, or service disruptions. Fault tolerance refers to the ability of a system to continue operating, even in the presence of faults. This is achieved by using various techniques that allow a system to detect, isolate, and recover from faults without affecting the system's overall functionality.

Fault tolerance is critical in systems that require high availability, reliability, and uninterrupted service, such as in aviation, banking, healthcare, and telecommunications.

2. Learning Outcomes

By the end of the lesson, students should be able to:

1. Demonstrate the understanding of the concept of fault tolerance and its significance in computing.
2. Identify different fault-tolerant schemes.
3. Explain the role of fault tolerance in critical systems like aviation and banking.

3. The Concept of Fault Tolerance

3.1 What is a Fault?

A **fault** is any abnormal condition that causes a system to deviate from the expected behavior. Faults can be categorized into different types:

- a. Transient Faults: Occur temporarily and disappear without any intervention.
- b. Intermittent Faults: Appear and disappear at irregular intervals.
- c. Permanent Faults: Persist until repaired or replaced.

3.2 Fault Tolerance

Fault tolerance is the property that enables a system to continue functioning even when one or more of its components fail. This is achieved through redundancy and other mechanisms that help in fault detection, isolation, and recovery.

3.3 Importance of Fault Tolerance in Modern Computing

- a. Ensures that systems remain operational with minimal downtime, which is critical for industries like healthcare and financial services.

- b. Improves the overall reliability of a system by preventing complete system failure due to component faults.
- c. In systems like aviation and autonomous vehicles, fault tolerance ensures that the system continues to operate safely, even in the event of a fault.

4. Fault-Tolerant Schemes


Fault-tolerant schemes are strategies used to detect, mask, and recover from faults. Some common fault tolerance schemes include:

4.1 Redundancy

Redundancy is one of the most widely used fault-tolerant schemes. It involves duplicating critical components or systems to ensure that a backup is available if the primary system fails. Redundancy can be applied in three domains as follows:

- a. *Hardware Redundancy*: Involves having multiple hardware components, such as processors, power supplies, or disks, that take over in case of failure.
- b. *Software Redundancy*: Uses diverse software versions or duplicate software modules that can substitute for one another in case of failure.
- c. *Information Redundancy*: Includes error-detecting and error-correcting codes like **Hamming Code** or **Reed-Solomon Code**, which ensure data integrity.

Diagram: Hardware Redundancy in a RAID System



An illustration showing multiple hard drives in a RAID setup, where redundant data is distributed across disks to ensure fault tolerance.


4.2 Replication

In replication, identical copies of processes or data are maintained. This ensures that if one instance fails, others can take over with no loss of data or service.

- *Data Replication*: Data is stored in multiple locations. If one data center fails, another can provide the same data.
- *Process Replication*: Critical processes are run on multiple machines or virtual environments. In case of failure in one, the others can continue executing without disruption.

Example of Process Replication in Distributed Systems

The diagram below shows a leader-based replication process. The Leader-based replication is an ideal choice for read-scaling scenarios where the read requests processed by a distributed system are far more than the number of write requests. This is often true of internet applications. The number of followers can be increased as the read load on the system increases.




An illustration of process replication across multiple servers in a distributed system, ensuring continuity of service.

4.3 Failover Systems

Failover systems automatically switch to a redundant or standby system when a fault is detected. This is common in critical services, such as web hosting and financial services, where downtime is unacceptable.

- **Cold Failover:** Involves switching to a backup system that is not running until a failure is detected. This approach has some delay due to startup time.
- **Hot Failover:** The standby system runs in parallel with the active system and can take over almost instantly when the primary system fails.

Diagram: Failover Mechanism



An illustration showing the primary system and a standby system, where the standby system takes over during a failure.


5. Role of Fault Tolerance in Critical Systems

Fault tolerance is essential in systems that require high reliability and availability. Some applications are found in banking, aviation, automobiles, defense etc.

5.1 In Aviation

In aviation, fault tolerance is critical for safety. Modern aircraft use **triple-modular redundancy (TMR)** in their flight control systems. In TMR, three independent systems run the same calculations. If one system gives a different result from the other two, it is automatically isolated, and the majority decision is taken as correct.

Diagram: Triple-Modular Redundancy in Aircraft




An illustration of TMR in a flight control system where three processors make parallel decisions.

5.2 In Banking

In banking, systems like **Automated Teller Machines (ATMs)** and **online transaction services** require continuous availability. Financial institutions use **data replication** across multiple data centers to ensure that transactions are processed even if one server fails. Additionally, **RAID systems** (Redundant Array of Independent Disks) are used to protect critical financial data against hardware failure.

Diagram: Data Replication in Banking Systems



An illustration showing how transaction data is replicated across multiple data centers to ensure availability.

5.3 In Medical Devices: In life-supporting medical devices like pacemakers and ventilators, fault tolerance ensures that the devices function reliably. Redundant components and error-correcting mechanisms help avoid failures that could result in life-threatening situations.

5.4 In Financial Systems: Banks, stock exchanges, and payment processing systems require fault tolerance to ensure continuous service availability and data integrity, even during hardware failures, network outages, or security breaches.

5.5 In Data Centers and Cloud Computing: Fault tolerance ensures that cloud services and data centers maintain high availability by employing redundant servers, storage systems, and failover mechanisms. This minimizes downtime and protects against data loss.

5.6 In Nuclear Power Plants: Fault tolerance is essential for the safe operation of nuclear power plants. It helps in managing and mitigating potential system failures that could lead to radiation leaks or meltdowns. Redundant sensors, control systems, and backup safety protocols are commonly used.

Fault tolerance is a fundamental aspect of designing modern computing systems, especially for critical applications like aviation, banking, and healthcare. By employing various schemes like redundancy, replication, and failover, fault-tolerant systems can continue to provide reliable and uninterrupted services, even in the face of faults and failures.

Evaluation Questions

1. **Define fault tolerance and explain its importance in modern computing systems.**
2. **Describe the following three fault-tolerant schemes with real-world examples.**
 - i. **Redundancy**
 - ii. **Failover Systems**
 - iii. **Replication**